

# CS 452: Kernel 1

Jason Sun (#20387090) and Shuo Chang (#20378235)

May 22, 2015

## 1 Program Operation

```
> load -b 0x00200000 -h 129.97.167.12 "ARM/sun-chang-team/k1.elf"
> go
```

All system calls required by assignment are supported:

`int Create(int priority, void (*code)())` Schedule a task with specified `priority` and function pointer `code`.

`int MyTid()` Return the task id for the calling task.

`int MyParentTid()` Return the task id of the parent of the calling task.

`void Pass()` No-op for entering the kernel.

`void Exit()` Exits the calling task and never schedule it again.

## 2 Kernel Details

### 2.1 Context Switch

From kernel space to user space:

1. Save user task `sp` into `r0`, variable register called `sp_`
2. Save user task `spsr` into `r1`, variable register called `spsr_`
3. Stack all kernel registers and return address `{r2-r12, lr}` on kernel stack.
4. Pop task `pc` as first word off `r0` (task `sp`) into kernel `lr` (restoring state for `movs`).
5. Restore `spsr` from `spsr_`.
6. Switch to system mode:
  - (a) Restore user task `sp` from `sp_`.

- (b) Unroll trap frame from task `sp` for registers `{r0-r12, lr}`. Note this is the task's `lr` and not the kernel's.
  - (c) Switch to supervisor mode `msr cpsr_c, #0xd3`.
7. Jump to userspace `movs pc, lr`.

**From user space to kernel:** Via `swi n` jumping to `KernelEnter` label

1. Switch to system mode:
  - (a) Store user task registers `{r0-r12, lr}` on user stack `pc`.
  - (b) Move `sp` into `r0`. Note this is `ssp_`.
2. Switch to supervisor mode.
3. Push `lr` (the task's `pc`) onto stack pointed by `r0` (task `sp`) and increment `r0`.
4. Move `spsr` into `r1`.
5. Pop kernel registers from kernel stack `{r2-r12, lr}`.
6. Restore `r0` (task `sp`) and
7. Restore `r1` (task `spsr`) into task descriptor.

### 2.1.1 Description in armlish

The piece of code responsible for context switch is:

```

register unsigned int *sp_ asm("r2") = active->sp; // r2 <- sp
register unsigned int spsr_ asm("r3") = active->spsr; // r3 <- spsr
*(sp_ + 1) = first.ret; // save ret on stack
unsigned int arg0, arg1;
asm volatile(

    "stmfd sp!, {r4-r12, lr}\n\t" // save kregs on kstack
    "ldmfd %0!, {lr}\n\t" // sp_ <- lr (the stored pc)
    "msr spsr, %1\n\t" // spsr <- spsr_
    "msr cpsr_c, #0xdf\n\t" // switch to system mode
    "mov sp, %0\n\t" // sp <- sp_
    "ldmfd sp!, {r0-r12, lr}\n\t" // pop task's registers
    "msr cpsr_c, #0xd3\n\t" // switch to supervisor mode
    "movs pc, lr\n\t" // jump to userspace

```

```

"KernelEnter:\n\t"           // label (swi jumps here)
"msr cpsr_c, #0xdf\n\t"      // switch to system mode
"stmfd sp!, {r0-r12, lr}\n\t" // store task registers
"mov %0, sp\n\t"             // sp_ <- sp save task's sp
"msr cpsr_c, #0xd3\n\t"      // switch to supervisor mode
"stmfd %0!, {lr}\n\t"         // sp + 0 <- lr save task's pc to stack
"mrs %1, spsr\n\t"           // spsr_ <- spsr save activity's spsr
"ldmfd sp!, {r4-r12, r14}\n\t" // unroll kregs from kstack

"mov %2, r0\n\t"             // copy arg0
"mov %3, r1\n\t"             // copy arg1

: "+r"(sp_), "+r"(spsr_), "=r"(arg0), "=r"(arg1) // output
: // input
: "r0", "r1" // force asmbldr not use any of these registers
);
active->sp = sp_;
active->spsr = spsr_;

```

### 2.1.2 Trap Frame

When the user does a syscall, a trap frame is set up and calls the kernel. The trap frame is pushed on the calling task's stack, storing its state (registers). The layout of registers stored is:

```

[      ] <-- SP after storing trap frame
[ PC   ]
[ LR   ]
[ r12  ]
[ r11  ]
[ ...  ]
[ r0   ] <-- SP at SWI instruction
[ ...  ]

```

Initializing the trap frame and returning from a call is written in assembly code, and can be found in `context_switch.s`, or inlined with `asm` operand. On return, the result of the syscall is stored in register `r0` and execution resumes at the point where the syscall occurred.

## 2.2 Syscalls

Syscalls defined in C functions, in `syscall.{c,h}` files. There is a `syscall` structure that contains the syscall type, and args 1 and args 2.

## 2.3 Tasks

A task can be created off a function pointer and represents a chunk of code to execute.

### 2.3.1 Task Descriptor

A TaskDescriptor struct holds:

- task id, contains an index into a global table of task descriptors pre-allocated.
- parent id, whoever called `Create()`.
- return value,
- stack pointer, `sp` and `spsr` are manipulated by the context switch.
- saved program status register
- and a pointer to the next task descriptor for singly linked list.

Currently only can create 128 tasks before failing to create more tasks.

### 2.3.2 Scheduling

Tasks each has a priority level. The scheduler tracks this tasks' priority via 32 ring buffer queues.

A bitmask keeps track of which of the 32 queues contains tasks. Using this bitmask, we efficiently computing the number of right leading zeroes in the bitmask with De Bruijn table lookup.

The kernel call `taskSchedule()` on each loop, and the queue with the highest priority is returned. The head of that queue is rotated to be the tail and the pointer is returned as the next task to be scheduled.

### 2.3.3 Task Creation

A task is created by specifying a priority, a function pointer, and parent task id. The `Create()` syscall is implemented by this function.

A task descriptor is filled in to the task table. Then a stack is allocated, a size of 4096 words. There is syscall to change a task's stack size. It also initializes a trap frame by setting `pc` to the value of the function pointer and saved stored program register.

Finally the kernel adds the task descriptor to priority queue.

### 2.3.4 Task Exit & Deletion

Once a task is removed from the priority queues, the task will not be scheduled again. No effort is made to reclaim task descriptors.

### 3 Source Code Location

Code is located under /u1/j53sun/cs452team/.

Compiling by running `make`, which also copies the local `kernel.elf` to /u/cs452/tftp/ARM/j53sun/cs452/.

File `md5sums`

8da586d949d31e239dfbe0c8356588f6	<code>bwio.c</code>
db0bab80ffcef52c8ce1a968c65587a9	<code>bwio.h</code>
d02b490ecfa9f94ca03ccb1004f23efe	<code>context_switch.h</code>
1ba8cd1b57c22116b57e96d22022cec5	<code>context_switch.s</code>
e9ecc0c507565cc766ec637a9aec3ab6	<code>cpsr.h</code>
6bf72fad920c3d9e326401a101e31ac0	<code>k1.lyx</code>
524bf9b87a0c352c3f62a46c15618ca3	<code>k1.pdf</code>
8fa76c583dac06f80d4028cdee20d7c3	<code>kernel.c</code>
e87799ad275ab3fd1199dba2ea334e5c	<code>linker.ld</code>
7c1b255735fd098a6ecd2b8a8903a0d9	<code>Makefile</code>
53fdea2ffa00ca6f9bbbea8f47d7b5ea	<code>readme.md</code>
d8e9046f472dc425d1cb3f884e0c939e	<code>scheduler.c</code>
0954bdc95abf5a2dfe291e04b8dedb80	<code>scheduler.h</code>
d6bdf5714a8d499da29f1064973580ae	<code>stdbool.h</code>
d8382c43b39efd0f066522396ddb41b5	<code>syscall.c</code>
702846c6401a9fddeb232c1244e03511	<code>syscall.h</code>
8e4e2530209d10a7f2a7b9c450cafd08	<code>task.c</code>
11d63ce61f395dde30e120fd11ed183a	<code>task.h</code>
2c5fc627ac5386f1f96a65d2f8dc9d67	<code>ts7200.h</code>
1e7dca0aa66b495b8d4f1e89490b88f1	<code>user_task.c</code>
962dc8dab4c71088e86d7c7c6bb9adc8	<code>user_task.h</code>

### 4 Program Output

```
RedBoot> go
Program completed with status 0
```

**Explanation**